

Dear Editors,

I submit “Verdict-style prompts induce moral over-resolution in large language models” as an Article in *Nature Machine Intelligence*.

When asked for a verdict on a contested case, large language models can make that case appear more settled than source-community votes and the model’s own distributional answer support. In a preregistered audit of five LLMs evaluated on SCRUPLES Anecdotes, I find that verdict-mode agreement estimates substantially exceeded both source-community vote support for the chosen label and the same model’s own distribution-mode probability for that same label, holding model, item and chosen label constant across formats. Repeated forced-choice outputs were similarly more concentrated than source-community votes. The pattern held across all five evaluated models and multiple robustness checks.

This matters for how LLM reliability is measured. A model that produces non-degenerate distributional outputs can still over-resolve disagreement when asked for a verdict; distribution-mode evaluation alone can miss it. The study contributes a within-model design for testing cross-format coherence that does not require perfect calibration to the source-community distribution, an approach applicable beyond the moral-judgment domain tested here.

The study is preregistered on OSF. Release-safe data, code and reproducibility materials are available through OSF, GitHub and Zenodo. No new human participants were recruited.

This manuscript is not under consideration at any other journal. The author declares no competing interests.

Yours sincerely,

Paul Briggs

Independent researcher, Los Angeles, CA, USA

[p.ivan.briggs@gmail.com](mailto:p.ivan.briggs@gmail.com)

ORCID: 0009-0002-9951-8661